



Innovate Summit 2017

Infrastructure as Code

Fredrik Skogman

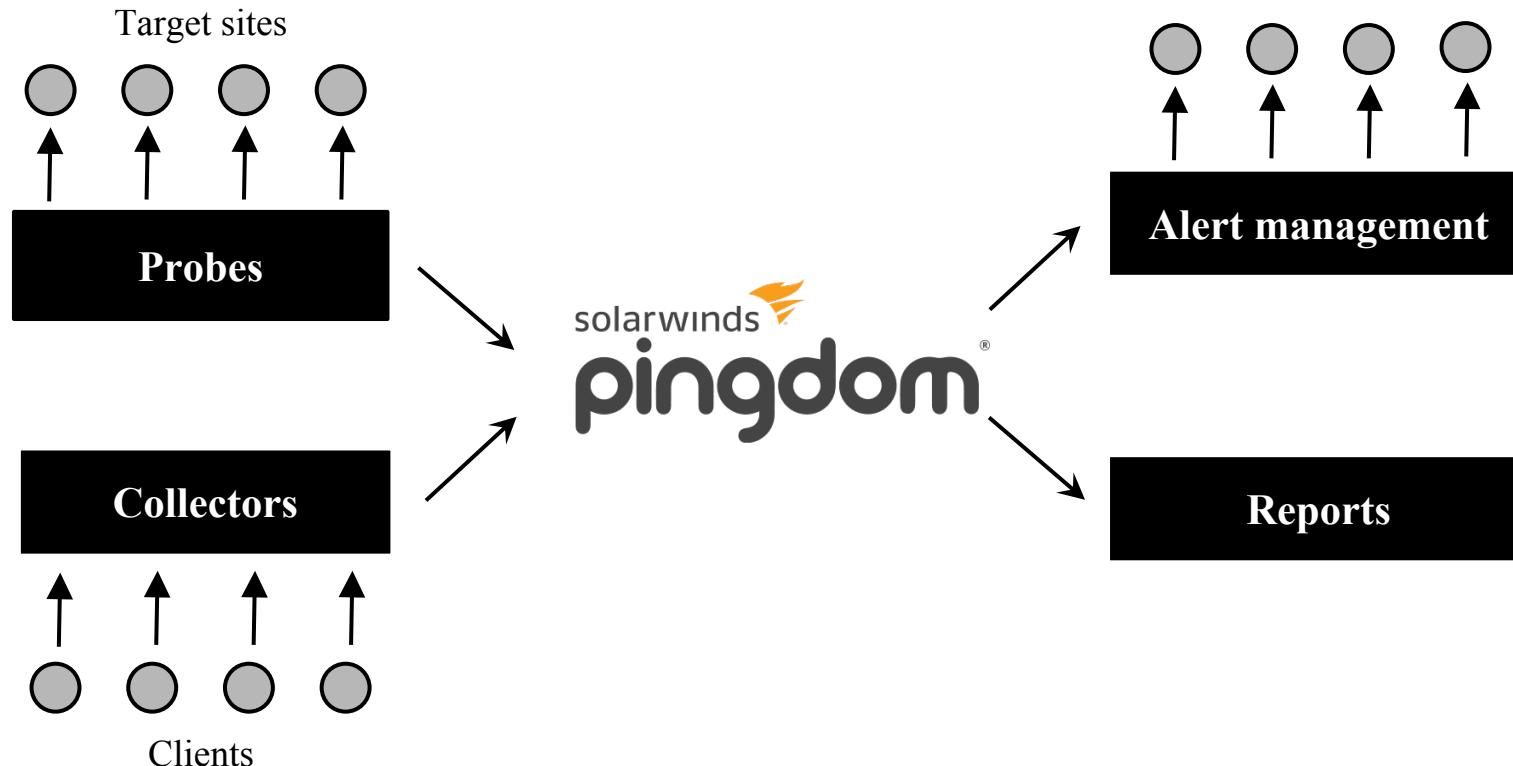


Fredrik Skogman

- Principal Architect at Pingdom, part of SolarWinds Cloud
- Located in Västerås, Sweden
- Primary focus is Pingdom product: back end, infrastructure, and front end



This is Pingdom



- Different services have different hardware requirements (instance type and count)
- Multiple database clusters with different hardware requirements
- Multiple event processing system (streaming data, message queues)
- Different services are deployed in different networks with different firewall rules
- Each service has its own set of permissions to access other services (databases, message queues, etc.)
- APIs located behind load balancers
- All logs are forwarded to a specific log management system
- Monitoring and alerting

- Manual management is bound to fail: human factor
 - Misconfiguration
 - Doing tasks in the wrong order, etc.
 - Non-existing or out-of-date documentation
- Lots of repetitive work
- Risk of snowflake servers/configuration drift
- Risk of misconfigured or missing monitoring/alerting



- Configuration management tools (Chef®, Puppet®, SaltStack®, etc.)
 - Usually used for server provisioning (installing patches, libraries, daemons, etc.)
 - Generation of configuration files
- Orchestration/provisioning tools (Terraform®, CloudFormation)
 - Usually used for hardware provisioning (instance type, number of instances, etc.)
 - Keeps a stored state of the system
 - Allows for trivial referencing across systems (e.g., map a database cluster to a set of application nodes)

- Declarative
 - Describes the desired state
 - Immutable infrastructure
- Imperative
 - Describes how to achieve the desired state
 - Can lead to config drift

Example (made up)

- launch_instances:

count: 10

image: ami-xyz

instance_type: hw.type.x

Example (made up)

```
resource "hw_instance" "example" {  
    count = 10  
  
    image = "ami-xyz"  
  
    instance_type = "hw.type.x"  
}
```

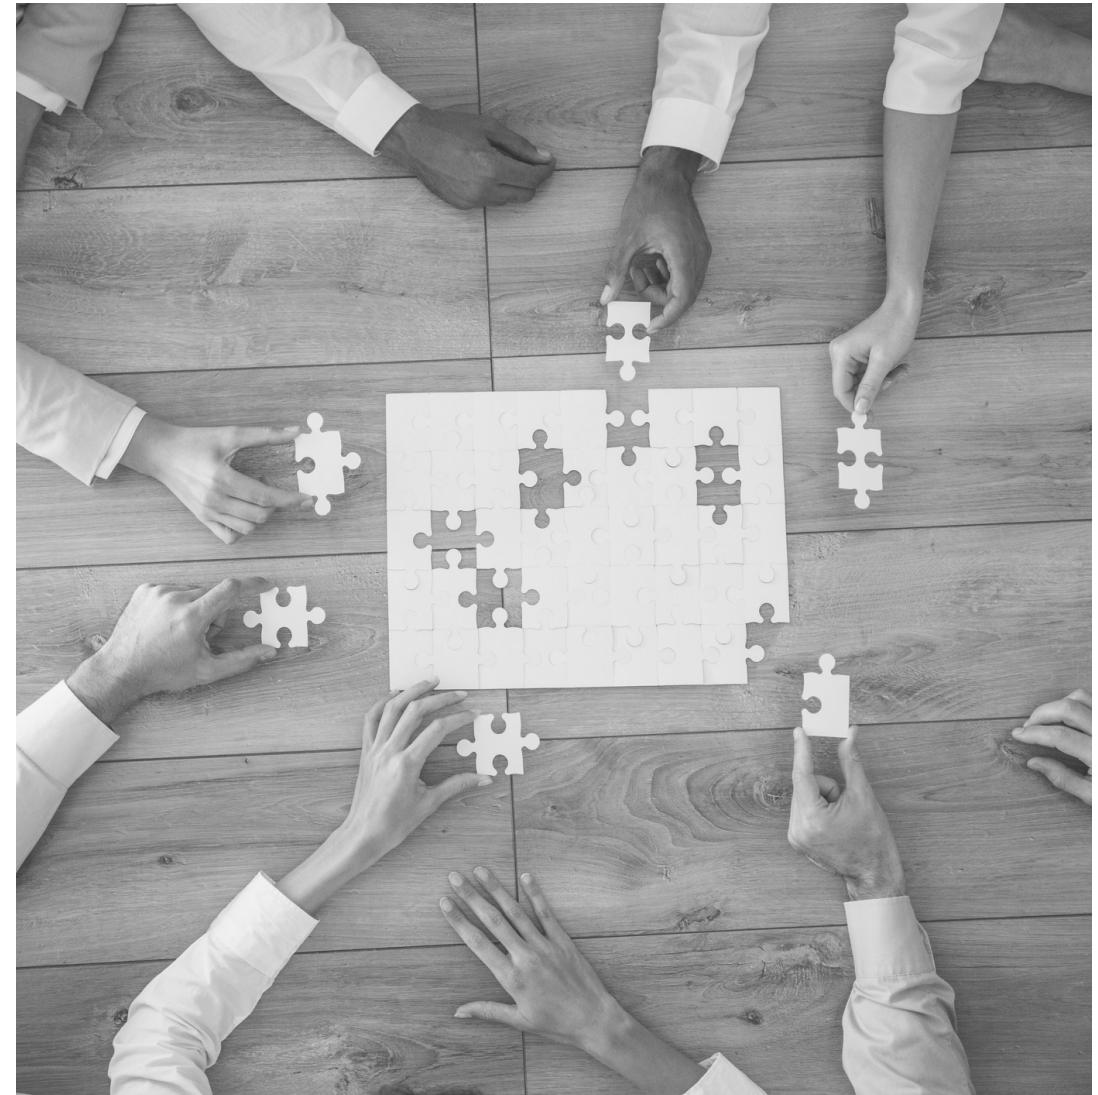
Declarative - continued

Example (Terraform)

```
resource "aws_security_group" "collector" {
    name      = "${var.pingdom_environment}-a-collector"
    vpc_id    = "${var.vpc}"
    ingress {
        from_port      = ${from_port}
        to_port        = ${to_port}
        protocol       = "TCP"
        security_groups = ["${aws_security_group.collector_elb.id}"]
    }
}

module "ecs_collector" {
    source          = ".../ecs"
    pingdom_environment = "${var.pingdom_environment}"
    load_balancers   = ["${aws_elb.collector_elb.name}"]
    num_instances    = "${var.collector_instance_count}"
    instance_type    = "${var.collector_instance_type}"
    security_groups  = ["${aws_security_group.collector.id}"]
    subnets          = ["${var.public_in_subnets}"]
}
```

- Easy (fast) to reproduce a system
- Version controlled infrastructure
- Safe rollbacks
- Easy to collaborate
 - Perform changes and create a PR
 - Share modules between teams
- Infrastructure can be unit tested
- Easy to manage multiple environments
- Safe updates to infrastructure – dry run

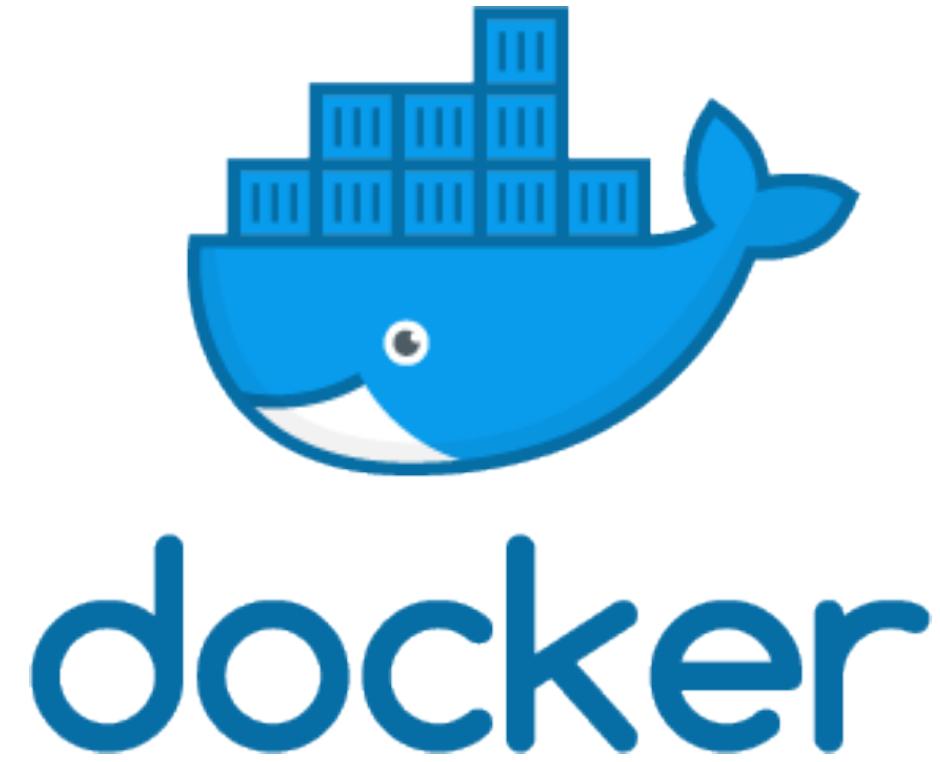


- Monitoring agents configured along with the infrastructure
- Alerting for both infrastructure and products
 - Different alert policies
 - Different alert targets



- SRE team owns the infrastructure
- Developers collaborate by creating PR when needed. SRE reviews
- Terraform for static resources
 - Reusable modules - one per service
 - Per environment overrides (e.g., hardware size and count)
- SaltStack for services (e.g., databases)
- Homogenous cluster for product deployments
 - Capture dependencies in containers

- Clear separation of modules and products
- Module is a self contained container
 - Containers are stateless
 - Configuration is added via environment variables
 - Scale up by deploying more container instances
- Each module has its own development life cycle
 - Peer review
 - Integration tests

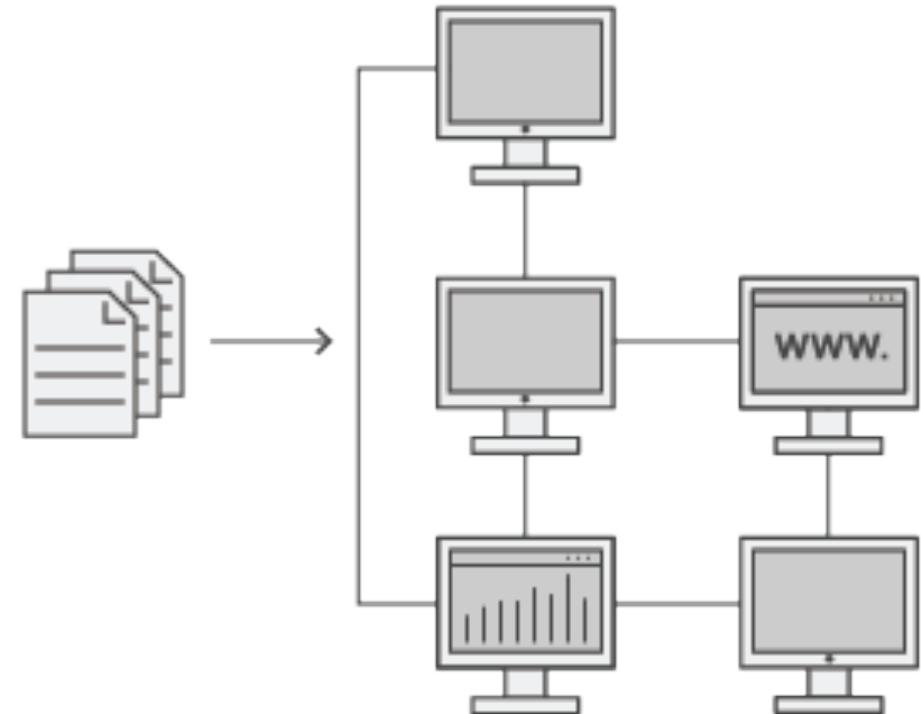


- A product is mainly metadata
 - What container to use?
 - Version of container
 - Configuration for each container
 - Credentials not part of configuration (added later)
- Products are version controlled
 - Peer review
 - Rollbacks, etc.
- Each product has its own development lifecycle

- Deployment is a single function with two arguments
 - Environment (production, system-test, etc.)
 - Product (including version)
- Deployment automatically performs the following tasks:
 - Choose hardware cluster to deploy to
 - Create a deployment artifact with necessarily credentials
 - Push artifact to chosen cluster
 - Spin up the product

- Per-product monitoring may need to be added by developers
 - Metrics
 - Use a log management system than can create alerts from logs
- APM (AppOptics™) can auto generate valuable metrics
- External monitoring to keep track of customers experience (Pingdom®)

- Infrastructure as application code
 - Agile
 - Iterative
 - Continuously improved infrastructure
- Dare to experiment—learn from failures
 - Rollback is easy, remember?



- Naming is important, namespace your components
- Stricter requirements for product development (such as 12-Factor App)
- Don't reinvent the wheel—use a mature CI & CD system
- Start small, think big



That's it!



Thank you!

The SolarWinds, SolarWinds & Design, Orion, and THWACK trademarks are the exclusive property of SolarWinds Worldwide, LLC or its affiliates, are registered with the U.S. Patent and Trademark Office, and may be registered or pending registration in other countries. All other SolarWinds trademarks, service marks, and logos may be common law marks or are registered or pending registration. All other trademarks mentioned herein are used for identification purposes only and are trademarks of (and may be registered trademarks) of their respective companies.